



Top 5 Deterministic Liar's Dice Strategies for a 5-Player Bot

In a 5-player, turn-based Liar's Dice tournament (no wilds, deterministic logic), the following five strategies have proven highly effective. Each strategy is described along with its core philosophy, pseudocode for integration into a WebWorker-based bot framework, and an evaluation of implementation complexity (within ~200ms per move, using deterministic logic).

1. Bayesian Inference Strategy

Core Logic & Philosophy: This strategy treats opponents' bids as informative signals about their hidden dice. Instead of assuming all unseen dice are uniformly random, the bot **updates its beliefs** whenever a player makes a bid. For example, if an opponent bids a high quantity of a certain face, it infers that opponent likely holds at least some of that face. By using Bayesian reasoning on the “**information exchanged by players' decisions**” ¹, the bot refines the probability that any given bid is truthful. It aims to call bluffs only when its updated model deems a claim very unlikely, and otherwise raise with bids that are plausible under the refined belief distribution. Over a 2500-game run, this bot can also adjust its prior assumptions based on aggregate outcomes (e.g. if bluffs are frequently caught, it may become more skeptical overall). The philosophy is to make **maximally informed decisions** by learning everything possible from the bidding history.

Pseudocode:

```
onmessage = (e) => {
  const { you, players, currentBid, history } = e.data.state;
  const myDice = you.dice;
  const totalDice = players.reduce((sum,p)=> sum + p.diceCount, 0);
  const unknownDiceCount = totalDice - myDice.length;

  // Belief model: initialize uniform probabilities for unseen dice
  // belief[playerId][face] = estimated count of 'face' dice that player holds.
  if (!self.belief) {
    self.belief = {};
    for (let p of players) {
      if (p.id !== you.id) {
        self.belief[p.id] = Array(7).fill(0);
        // Start with expected count = (p.diceCount * 1/6) for each face
        for (let face = 1; face <= 6; face++) {
          self.belief[p.id][face] = p.diceCount / 6;
        }
      }
    }
  }
}
```

```

        }

    }

    // Update beliefs based on the latest bid in history (if any new info)
    if (history.length > 0) {
        const lastAction = history[history.length - 1];
        if (lastAction.action === 'raise') {
            const bidder = lastAction.actor;
            const { quantity: q, face: f } = lastAction;
            // If an opponent raised on face f, assume they likely have at least one
            f.
            if (bidder !== you.id) {
                // Increase belief that this bidder holds face f (up to their dice
                count)
                self.belief[bidder][f] =
                    Math.min(players.find(p=>p.id==bidder).diceCount,
                            Math.max(self.belief[bidder][f], 1));
            }
        }
    }

    // Helper: calculate probability current bid (qty of face) is true under
    beliefs
    function probabilityBidTrue(qty, face) {
        let have = myDice.filter(d => d === face).length;
        let need = qty - have;
        if (need < 0) return 1; // I alone exceed the bid
        // Estimate expected count of `face` among unknown dice using beliefs
        let expectedUnknown = 0;
        for (let p of players) {
            if (p.id !== you.id) {
                // use belief or default 1/6 expectation for this face
                expectedUnknown += self.belief[p.id]
                    ? Math.min(p.diceCount, self.belief[p.id][face])
                    : p.diceCount/6;
            }
        }
        // Use a binomial-tail approximation around the expectedUnknown
        const mean = expectedUnknown;
        const variance = expectedUnknown * (5/6); // rough variance assuming binary
        outcome per die
        const stddev = Math.sqrt(variance);
        // Approximate probability at least `need` successes (Chebyshev bound or
        normal approx)
        let prob = 0;
        if (need <= mean) {
            prob = 1; // if need is below expected, assume highly likely
        } else {

```

```

        // if need > mean, estimate probability with tail of normal distribution
        const z = (need - mean) / (stddev || 1);
        // simple normal tail approximation
        prob = Math.exp(-0.5 * z * z);
    }
    return Math.max(0, Math.min(1, prob));
}

if (!currentBid) {
    // Opening bid: choose face I have the most of, quantity near expectation
    const counts = Array(7).fill(0);
    myDice.forEach(d => counts[d]++;
    let bestFace = counts.indexOf(Math.max(...counts));
    if (bestFace < 1) bestFace = 1;
    // expected total of bestFace = my count + expected others
    let expTotal = counts[bestFace] + (unknownDiceCount * (1/6));
    let openQty = Math.max(1, Math.floor(expTotal));
    postMessage({ action: 'raise', quantity: openQty, face: bestFace });
    return;
}

const { quantity: currQ, face: currF } = currentBid;
const claimProb = probabilityBidTrue(currQ, currF);

// Decide to call liar or raise
const LIAR_THRESH = 0.20; // base threshold for calling liar
// If our belief-adjusted probability of truth is very low, call liar
if (claimProb < LIAR_THRESH) {
    postMessage({ action: 'liar' });
    return;
}

// Otherwise, attempt a minimal credible raise
// Find the smallest legal raise (either +1 quantity or higher face)
const raiseOptions = [{q: currQ + 1, f: currF}].concat(
    currF < 6 ? [{q: currQ, f: currF + 1}] : []);
for (let opt of raiseOptions) {
    if (probabilityBidTrue(opt.q, opt.f) >= LIAR_THRESH) {
        postMessage({ action: 'raise', quantity: opt.q, face: opt.f });
        return;
    }
}

// If no safe raise, just call liar as last resort
postMessage({ action: 'liar' });
}

```

Implementation Notes:

- *Complexity*: High – maintaining and updating a belief distribution for each opponent adds complexity. The pseudocode above uses a simplified heuristic belief update rather than a full Bayesian computation over all dice combinations (which would be intractable for 5 players). Despite simplifications, the logic is more involved than a basic threshold strategy.
- *Data Structures*: A persistent `belief` object (stored on `self`) holds an estimated count of each face for each opponent. This requires memory proportional to $O(\text{numPlayers} * \text{faces})$, which is small (e.g. 5 players * 6 faces). More sophisticated Bayesian updates could use dynamic programming or particle filtering, but those would be more complex.
- *Computational Cost*: Moderate. The example uses an approximate probability calculation that runs in $O(\text{numPlayers})$ per decision. A more exact approach might iterate over possible distributions or run simulations, which could push the time budget. Tuning the number of belief updates or samples is necessary to stay within 200ms.
- *Determinism*: Fully deterministic. The strategy uses no random calls – it derives decisions from the current game state and history. (If Monte Carlo sampling were used for Bayesian updates, it would use the provided seeded `Math.random` to remain reproducible.) Given identical game states, it will always produce the same action.

2. Monte Carlo Simulation Planner

Core Logic & Philosophy: This strategy uses **Monte Carlo simulation** to evaluate moves. Rather than rely purely on analytical probability, the bot **simulates many possible dice distributions and outcomes** to decide whether to raise or call. Essentially, the bot “rolls out” the scenario hundreds or thousands of times in its head: it fills in all unknown dice randomly (according to known dice counts) and then simulates what would happen if it calls the bluff now versus if it raises the bid. By comparing the frequency of successful outcomes, the bot chooses the action with the best empirical success rate. This approach can incorporate more complex dynamics (e.g., potential reactions of other players) if a model for opponents’ behavior is assumed during simulation. The philosophy is to approximate a mini lookahead search in this imperfect-information game by brute-force sampling, which captures nuances that straightforward probability calculations might miss (such as the distribution of outcomes when the round continues). The result is a **data-driven, tuning-free decision** that tends to align with optimal play over many trials.

Pseudocode:

```
onmessage = (e) => {
  const { you, players, currentBid } = e.data.state;
  const myDice = you.dice;
  const totalDice = players.reduce((sum,p)=> sum + p.diceCount, 0);
  const unknownDiceCount = totalDice - myDice.length;

  // Helper: simulate a random assignment of all unknown dice
  function randomUnknownDice() {
    // Create an array representing all unknown dice (not including mine)
    const unknown = [];
    for (let p of players) {
      if (p.id !== you.id) {
```

```

    // assign random faces for each of this player's dice
    for (let i = 0; i < p.diceCount; i++) {
        // Using deterministic seeded Math.random provided by framework
        const face = Math.floor(Math.random() * 6) + 1;
        unknown.push(face);
    }
}
return unknown;
}

// Count occurrences of a face in an array of dice
function countFace(diceArray, face) {
    return diceArray.filter(d => d === face).length;
}

// Decide an opening bid if no bid yet (choose a reasonable starting bid)
if (!currentBid) {
    // Open with my highest-count face at roughly average quantity
    const counts = Array(7).fill(0);
    myDice.forEach(d => counts[d]++;
    let bestFace = counts.indexOf(Math.max(...counts));
    if (bestFace < 1) bestFace = 1;
    const expUnknown = unknownDiceCount / 6; // expected count of bestFace in
unknown
    const openQty = Math.max(1, Math.round(counts[bestFace] + expUnknown));
    postMessage({ action: 'raise', quantity: openQty, face: bestFace });
    return;
}

const { quantity: currQ, face: currF } = currentBid;

// Simulation parameters
const N = 1000; // number of simulations (tune based on 200ms budget)
let liarWins = 0, liarTotal = 0;
let raiseWins = 0, raiseTotal = 0;

for (let t = 0; t < N; t++) {
    const unknown = randomUnknownDice();
    // Current bid outcome if we call liar:
    liarTotal++;
    const actualCount = countFace(unknown.concat(myDice), currF);
    const claimTrue = (actualCount >= currQ);
    if (!claimTrue) {
        // Claim was false, calling liar would win (opponent loses a die)
        liarWins++;
    }
    // else claim was true, calling liar would lose (we lose a die)
}

```

```

    // Outcome if we raise minimally (quantity+1 of same face):
    // We assume that if our new bid is false, eventually someone will call and
    we'll lose;
    // if it's true, either it goes through or a liar call fails, meaning we
    survive or someone else loses.
    raiseTotal++;
    const newBidQ = currQ + 1;
    const newBidF = currF;
    const actualCountNew = countFace(unknown.concat(myDice), newBidF);
    const newClaimTrue = (actualCountNew >= newBidQ);
    if (newClaimTrue) {
        // Our raise could be upheld (good outcome for us, not losing a die in
        this round)
        raiseWins++;
    }
    // if false, assume we'll eventually be caught and lose (so no increment to
    raiseWins)
}

// Estimate success probabilities
const p_callWin = liarWins / Math.max(1, liarTotal);      // probability we win
by calling liar
const p_raiseWin = raiseWins / Math.max(1,
raiseTotal); // probability we come out safe by raising

// Decision: choose action with higher chance to avoid losing a die
if (p_callWin > p_raiseWin) {
    postMessage({ action: 'liar' });
} else {
    // Here we choose the minimal raise; could also simulate other raise options
    similarly
    postMessage({ action: 'raise', quantity: currQ + 1, face: currF });
}
};

```

Implementation Notes:

- *Complexity:* Moderate to High. The core logic (random sampling of dice outcomes) is straightforward, but ensuring the simulation is fast enough is the challenge. The pseudocode uses $N = 1000$ trials as an example; this may need adjustment to fit in <200ms depending on the environment's speed. Each trial involves random number generation and counting – relatively cheap operations – so a few thousand iterations are plausible within the time budget.
- *Data Structures:* The simulation uses simple arrays for unknown dice and counters. No large data structures are needed, but storing intermediate results for reuse (e.g., precomputing one random allocation and tweaking it) could optimize speed.
- *Computational Cost:* Tuning is required. For example, 1000 simulations per move might take on the order of tens of milliseconds in JavaScript, which is acceptable. The cost scales linearly with the number of samples

`N`. We must also simulate *both* calling and raising scenarios. To keep it deterministic, we use the seeded `Math.random` provided by the WebWorker (which is consistent across runs). If performance is an issue, `N` can be reduced or made adaptive (do more samples only when a decision is very close).

- **Determinism:** The strategy is stochastic in methodology but **deterministic in execution**. It relies on `Math.random` which has been seeded for the tournament run `②`, meaning the sequence of "random" outcomes is fixed given the initial seed. Thus, the bot will make the same decisions every time for identical game states. There is no hidden internal state except the progression of the RNG, which is globally deterministic. One consideration: because the RNG state persists, the exact sequence of random draws might depend on how many simulations were run in earlier turns. To maintain consistency, the number of samples `N` should be fixed or solely dependent on state (not, say, dynamically on time remaining). With careful implementation, the Monte Carlo planner yields a reproducible policy that approximates ideal play through large-sample evaluation.

3. Opponent Modeling & Exploitation

Core Logic & Philosophy: This strategy focuses on **reading the opponents** and adapting to their play styles `③`. The bot tracks each opponent's behavior over the course of a game (using the per-round history, since seat IDs are anonymized each new match). It gathers statistics such as how often each player raises versus calls "**LIAR**", the typical magnitude of their bids, and whether their bluffs succeed or fail. Using this data, the bot categorizes opponents into profiles – e.g. *conservative* (rarely bluff, often call), *aggressive bluffer* (frequently make large raises), *cautious* (only bid when very safe), *calling-station* (calls bluffs at the slightest doubt), etc. The bot then **exploits these tendencies**. For example, if the previous bidder is known to bluff often, this bot will call them out with less hesitation (i.e. require a higher probability that the claim is true before not calling liar). Conversely, if an opponent has been very truthful (never caught bluffing), the bot gives them the benefit of the doubt and won't call liar unless the claim is extremely unlikely. When the bot itself is raising, it considers the next player in turn: if the next player is very timid about calling, our bot can safely issue a bolder (more bluff-like) raise knowing that player will likely pass the turn; if the next player is very call-happy, our bot will only make raises that it can strongly back up with its own dice. This human-inspired strategy mirrors how skilled players "**watch for patterns to mix up their bids**" and calls `③` – effectively playing the opponents, not just the odds.

Pseudocode:

```
onmessage = (e) => {
  const { you, players, currentBid, history } = e.data.state;
  const myDice = you.dice;
  const totalDice = players.reduce((sum,p)=> sum + p.diceCount, 0);

  // Persistent memory for opponent stats within a match
  if (!self.stats) self.stats = {}; // { [playerId]: { raises: 0, liars: 0,
  totalActions: 0, bluffFails: 0 } }

  // Update stats from history (e.g., look at the last action)
  if (history.length) {
    const last = history[history.length - 1];
    if (!self.stats[last.actor]) {
```

```

        self.stats[last.actor] = { raises: 0, liars: 0, totalActions: 0,
bluffFails: 0 };
    }
    if (last.action === 'raise') {
        self.stats[last.actor].raises += 1;
        self.stats[last.actor].totalActions += 1;
    } else if (last.action === 'liar') {
        self.stats[last.actor].liars += 1;
        self.stats[last.actor].totalActions += 1;
        // If liar was called, we can detect if it was a correct call by looking
        at the next history entry
        // (In actual implementation, we might get an outcome event. Here assume
        history logs outcome of liar calls.)
        // Pseudocode assumes an outcome record or some way to increment
        bluffFails for the bidder if bluff caught.
    }
    // Note: We could also update bluffFails by checking if a raise resulted in
    a loss of die, etc., if that info is available.
}

// Define dynamic thresholds based on opponent styles
const baseLiarThreshold = 0.22; // baseline probability to call liar
let liarThreshold = baseLiarThreshold;
let raiseThreshold = 0.40; // baseline probability needed to raise
credibly

// Identify the opponent who made the current bid (if any) and the next
opponent in turn order
let prevBidderId = null;
if (currentBid && history.length) {
    // last raise in history is the current bid
    const lastRaise = history.slice().reverse().find(h => h.action === 'raise');
    prevBidderId = lastRaise ? lastRaise.actor : null;
}
// Find next player in turn (assuming players array is in turn order and you
are at your index)
let myIndex = players.findIndex(p => p.id === you.id);
let nextPlayerId = players[(myIndex + 1) % players.length].id;
if (nextPlayerId === you.id) {
    // if I'm last, nextPlayerId is myself (wrap-around), skip self
    nextPlayerId = players[(myIndex + 2) % players.length]?.id;
}

// Adjust liarThreshold based on previous bidder's behavior
if (prevBidderId && self.stats[prevBidderId]) {
    const oppStats = self.stats[prevBidderId];
    const raiseRate = oppStats.raises / Math.max(1, oppStats.totalActions);
    const liarRate = oppStats.liars / Math.max(1, oppStats.totalActions);
}

```

```

    // If opponent raises a lot relative to calling (aggressive bluffer), be
    more skeptical
    if (raiseRate > 0.7) {
        liarThreshold = 0.30; // call liar more often (require only <30% chance
        to call liar)
    }
    // If opponent rarely bluffs (low raiseRate or high liarRate), be more
    lenient (harder to call them liar)
    if (liarRate > 0.5) {
        liarThreshold =
        0.15; // they call others often -> likely honest, so need <15% chance to call
        them
    }
}

// Adjust raiseThreshold based on next player's likely response
if (self.stats[nextPlayerId]) {
    const nextStats = self.stats[nextPlayerId];
    const nextCallRate = nextStats.liars / Math.max(1, nextStats.totalActions);
    // If next player calls liar frequently (aggressive caller), only raise with
    strong confidence
    if (nextCallRate > 0.5) {
        raiseThreshold = 0.60; // need ≥60% chance our raise is true to attempt
        it
    }
    // If next player rarely calls liar (timid), we can bluff more (lower
    threshold)
    const nextRaiseRate = nextStats.raises / Math.max(1,
    nextStats.totalActions);
    if (nextCallRate < 0.2 && nextRaiseRate < 0.5) {
        raiseThreshold = 0.25;
    }
}

// Probability helper (basic like baseline)
function probabilityAtLeast(face, qty) {
    // known count of face in my dice
    const myCount = myDice.filter(d => d === face).length;
    const need = Math.max(0, qty - myCount);
    const p = 1/6;
    // Binomial tail for need successes in unknownDiceCount trials
    // (Using a simple approximation or small loop since unknownDiceCount is
    manageable)
    let prob = 0;
    for (let k = need; k <= unknownDiceCount; k++) {
        // (Compute C(unknownDiceCount, k)*p^k*(1-p)^(unknownDiceCount-k))
        // ...omitted for brevity...
    }
}

```

```

        return prob;
    }

    if (!currentBid) {
        // Opening bid: pick my highest-count face with quantity = that count or 1
        const counts = Array(7).fill(0);
        myDice.forEach(d => counts[d]++;
        let bestFace = 1;
        for (let f = 1; f <= 6; f++) {
            if (counts[f] > counts[bestFace]) bestFace = f;
        }
        const qty = Math.max(1, counts[bestFace]);
        postMessage({ action: 'raise', quantity: qty, face: bestFace });
        return;
    }

    const { quantity: currQ, face: currF } = currentBid;
    const claimProb = probabilityAtLeast(currF, currQ);

    // Decide to call liar based on adjusted threshold
    if (claimProb < liarThreshold) {
        postMessage({ action: 'liar' });
        return;
    }

    // Try to find a raise that meets our confidence threshold
    const legalRaises = [{q: currQ+1, f: currF}].concat(
        currF < 6 ? [{q: currQ, f: currF+1}] : []);
    for (let r of legalRaises) {
        if (probabilityAtLeast(r.f, r.q) >= raiseThreshold) {
            postMessage({ action: 'raise', quantity: r.q, face: r.f });
            return;
        }
    }

    // If no safe raise, just call liar if current claim is borderline (or make
    // minimal raise otherwise)
    if (claimProb < 0.35) {
        postMessage({ action: 'liar' });
    } else {
        // minimal bump to keep game moving
        postMessage({ action: 'raise', quantity: currQ+1, face: currF });
    }
}

```

Implementation Notes:

- *Complexity:* Moderate. Tracking opponent behavior requires bookkeeping but not heavy computation. The

pseudocode maintains a `stats` object and updates it each turn; this is simple counting and conditional logic. Profiling opponents is conceptually straightforward, though choosing good thresholds for "aggressive" vs "conservative" is a tuning task (the example uses arbitrary 70%/50% cutoffs for illustration).

- *Data Structures*: A dictionary of opponent statistics, keyed by anonymized player ID for the current game. This uses minimal memory (a few counters per opponent). The game's `history` provides the needed data. Note that because seat IDs change each match, the bot cannot carry these profiles across games – it effectively "re-learns" each game, which is within the rules (no cross-game identity tracking).
- *Computational Cost*: Low. Updating counts and adjusting thresholds are $O(1)$ per turn. The probability calculation uses a binomial formula similar to the baseline (could reuse baseline's functions), which is efficient for up to ~25 unknown dice. Overall, the strategy easily operates well under 200ms.
- *Determinism*: Deterministic. The logic consists of fixed rules and arithmetic on game state data. There is no random component. Even though it adapts to opponents, it does so in a fixed way given the observed history – i.e. two runs with the same sequence of actions will produce the same adjustments. One subtlety is that this strategy may create a **feedback loop** in long matches: if it identifies someone as a bluffer and calls them, that opponent's stats will shift. However, since all adjustments are based on concrete past actions, the behavior remains consistent and replayable. By mimicking how humans read opponent tendencies, this bot can exploit predictable players and respond optimally to different styles.

4. Dice-Count Adaptive Risk Strategy

Core Logic & Philosophy: This strategy adapts its level of risk-taking based on **the dice counts (or "stacks") of all players** and the stage of the game. The core idea is to dynamically shift between conservative and aggressive play depending on one's advantage. When the bot has **more dice than most opponents (a large stack)**, it plays more conservatively: it avoids unnecessary liar calls and only bluffs with bids that it can reasonably support. The reasoning is that with a lead, it's better to let opponents knock each other out rather than risk losing dice on a bad call. On the flip side, if the bot is **down to a few dice (disadvantaged)**, it becomes more willing to take risks – calling borderline bluffs or making daring raises – because playing safe when behind often leads to elimination anyway. This aligns with human tournament intuition: big leaders protect their lead, while underdogs gamble to get back in the game. Moreover, the strategy accounts for the **overall number of dice remaining in the game**. In early rounds with many dice in play, moderate bluffs are statistically more likely to succeed (more dice = more uncertainty) ⁴. In late-game scenarios with few dice, wild claims are usually obvious lies (since "the risk of bluffing increases with fewer dice remaining" ⁵), so the bot tightens up significantly. By adjusting its liar-call threshold and bidding aggression according to these factors, this strategy maintains an optimal risk-reward balance throughout the match.

Pseudocode:

```
onmessage = (e) => {
  const { you, players, currentBid } = e.data.state;
  const myDiceCount = you.dice.length;
  const totalDice = players.reduce((sum,p)=> sum + p.diceCount, 0);
  const unknownDiceCount = totalDice - myDiceCount;

  // Determine relative rank by dice count
  const sortedCounts = players.map(p => p.diceCount).sort((a,b)=>b-a);
```

```

const myRank = sortedCounts.indexOf(myDiceCount) + 1; // 1 = highest dice
count
const maxDice = sortedCounts[0];
const minDice = sortedCounts[sortedCounts.length - 1];

// Base thresholds
let liarThreshold = 0.20; // base probability below which to call liar
let raiseThreshold = 0.40; // base probability needed to raise

// Adjust thresholds based on stack rank
if (myDiceCount === maxDice) {
    // I'm one of the leaders: be more cautious (harder to call liar, safer
raises)
    liarThreshold = 0.15; // require claim <15% likely to call liar
    raiseThreshold = 0.50; // need ≥50% chance for a raise to be made
}
if (myDiceCount === minDice) {
    // I'm at the bottom: be more aggressive (call more, bluff more)
    liarThreshold = 0.30; // call if claim <30% likely
    raiseThreshold = 0.30; // only need 30% chance to attempt a raise
}

// Adjust for overall game stage (fewer dice => be more truthful and
selective)
if (totalDice < 10) { // late game scenario, very few dice left
    // Reduce bluffing because outcomes are more certain with fewer dice
    raiseThreshold = Math.max(raiseThreshold, 0.50);
    // Possibly increase liar calls slightly since big claims are easier to
catch
    liarThreshold = Math.max(liarThreshold, 0.25);
}

// Probability function as in baseline
function probabilityAtLeast(face, qty) {
    let have = 0;
    for (let d of (you.dice||[])) {
        if (d === face) have++;
    }
    const need = Math.max(0, qty - have);
    const p = 1/6;
    // simple binomial tail (could call baseline's function)
    let prob = 0;
    for (let k = need; k <= unknownDiceCount; k++) {
        // compute probability of exactly k occurrences of 'face' in unknown dice
        // and sum for k>=need.
        // ... (omitted for brevity; assume this calculation exists) ...
    }
    return prob;
}

```

```

}

if (!currentBid) {
  // Opening bid: conservative if leading, aggressive if trailing
  // Choose face I have most of. Quantity: if I'm trailing, round up
  expectation; if leading, round down.
  const myFaceCounts = Array(7).fill(0);
  for (let d of you.dice) myFaceCounts[d]++;
  let bestFace = 1, bestCount = -1;
  for (let f = 1; f <= 6; f++) {
    if (myFaceCounts[f] > bestCount) { bestFace = f; bestCount =
    myFaceCounts[f]; }
  }
  const expectedUnknown = unknownDiceCount * (1/6);
  let qty = bestCount;
  if (myDiceCount === minDice) {
    // behind - be a bit bolder
    qty = Math.max(1, Math.ceil(bestCount + expectedUnknown));
  } else {
    // ahead or normal - be a bit conservative
    qty = Math.max(1, Math.floor(bestCount + expectedUnknown) - 1);
  }
  postMessage({ action: 'raise', quantity: qty, face: bestFace });
  return;
}

const { quantity: currQ, face: currF } = currentBid;
const claimProb = probabilityAtLeast(currF, currQ);

// Liar call decision
if (claimProb < liarThreshold) {
  postMessage({ action: 'liar' });
  return;
}

// Raise decision: find a minimal raise meeting raiseThreshold
const options = [{ q: currQ+1, f: currF }];
if (currF < 6) options.push({ q: currQ, f: currF+1 });
for (let opt of options) {
  if (probabilityAtLeast(opt.f, opt.q) >= raiseThreshold) {
    postMessage({ action: 'raise', quantity: opt.q, face: opt.f });
    return;
  }
}

// If no raise is safe and we also aren't confident enough to call liar, nudge
up minimally

```

```

    postMessage({ action: 'raise', quantity: currQ+1, face: currF });
}

```

Implementation Notes:

- *Complexity*: Low. The strategy mainly tweaks a couple of threshold values based on simple metrics (dice counts). The conditional logic for adjusting thresholds is straightforward. This makes the code easy to implement and verify.
- *Data Structures*: It uses basic variables and computations on the `players` list (which contains each player's remaining dice count). Sorting the dice counts to find `maxDice` and `minDice` is $O(n \log n)$ (with $n=5$ players, negligible). Alternatively, one can scan for max/min in $O(n)$. No persistent memory across rounds is needed except perhaps a counter for rounds if further stage nuance is desired.
- *Computational Cost*: Very small. Even with the probability calculation (binomial summation) for each potential action, the overhead is minor given at most 25 unknown dice and a couple of raise options. The adjustments to thresholds and decision logic add only a few arithmetic operations and comparisons. Running well under 200ms is virtually guaranteed.
- *Determinism*: Completely deterministic. The strategy doesn't rely on any random choices. Given the same configuration of dice counts and current bid, it will always produce the same decision. One thing to note is that this strategy can create *predictable patterns* (e.g., a leading bot almost never calls unless extremely safe, a trailing bot calls more often). While this determinism might be exploitable in theory (opponents could adjust knowing your behavior by dice count), the tournament's anonymity and simultaneous adaptation by others make it hard to target. In practice, the dice-count adaptive strategy has been effective at preserving leads and giving underdogs a fighting chance by **calibrating bluffing risk to the game state**, reflecting sound tournament tactics.

5. Game-Theoretic Equilibrium Strategy (CFR/NFSP-Based)

Core Logic & Philosophy: This strategy is derived from **game-theoretic analysis and AI self-play training**, aiming to approximate the *Nash equilibrium* strategy for Liar's Dice. Instead of hand-crafted rules, it leverages an offline-computed policy (e.g., via **Counterfactual Regret Minimization (CFR)** or **Neural Fictitious Self-Play (NFSP)**) that tells the bot how to act optimally in any given state ⁶ ⁷. The core philosophy is to minimize exploitability: it balances bluffing and truthful play in such a way that no opponent can easily take advantage. In practice, an equilibrium strategy will sometimes make surprising moves (like an occasional bluff with a weak hand or an unexpected pass with a strong hand) to avoid revealing private information ⁸. However, since our environment requires determinism, this bot implements a *deterministic approximation* of the mixed-strategy equilibrium. For example, it might always choose the highest-probability action recommended by the equilibrium strategy (rather than randomizing). The result is a highly sophisticated policy that has learned, through simulation of millions of rounds, the best responses to any situation. Such a bot can adjust to the table implicitly – *for instance, it learns when to bluff less because opponents call often, or vice versa, as these patterns are encoded in its strategy table*. In essence, this approach attempts to play *perfect Liar's Dice* (within the no-randomness constraint), making it extremely strong and resilient. Notably, research has shown that in multi-player Liar's Dice variants, neural self-play agents can achieve **nearly unexploitable performance** ⁷ ⁹, which this strategy emulates.

Pseudocode:

```

// Pseudocode: Using a precomputed strategy (policy) derived from equilibrium
analysis
// For illustration, assume we have a function getEquilibriumAction(state) ->
{action, quantity, face}
// which encapsulates the bot's trained strategy. This could be a lookup table
or neural network.

onmessage = (e) => {
  const state = e.data.state;
  // The state includes you.dice, players (dice counts), currentBid, and
  history.
  // Construct a simplified state representation that indexes into our policy
  table.
  const myDiceSorted = [...state.you.dice].sort().join(''); // e.g. "146" if
  dice are [1,4,6]
  const currentBidStr = state.currentBid
    ? `${state.currentBid.quantity}${state.currentBid.face}`
  `

    : "none";
  // We might also include some recent actions or aggregated info if needed.

  // Compose an information-set key (this is a simplified key; real strategy
  might use more detail):
  const infoKey = `${myDiceSorted}|${currentBidStr}|$`+
  {state.players.map(p=>p.diceCount).join(',')}`;

  // Lookup the recommended action from the equilibrium policy map (precomputed
  offline)
  let actionRec;
  if (self.equilibriumPolicy && self.equilibriumPolicy[infoKey]) {
    actionRec = self.equilibriumPolicy[infoKey];
    // example value: {action: 'raise', quantity: 5, face: 3, mixProb: {raise:
    0.7, liar:0.3}}
  } else {
    // Fallback: if state not in policy (unlikely if policy is comprehensive),
    use a default safe strategy
    actionRec = { action: state.currentBid ? 'liar' : 'raise', quantity:
    state.currentBid ? null : 1, face: state.currentBid ? null : 1 };
  }

  // Deterministic choice: pick the action with highest probability in the mixed
  strategy (or a fixed rule)
  // mixProb might suggest randomness, but we choose deterministically for
  reproducibility.
  // E.g., if policy says 70% raise 5x3, 30% call, we'll always take the raise
  (since it's highest probability).
  let finalAction;

```

```

if (actionRec.mixProb) {
    // Choose argmax of mixProb (deterministic choice of the most favored
    action)
    let bestAct = actionRec.action;
    // (In practice, we might just assume actionRec.action is already the best
    deterministic choice.)
    finalAction = { action: bestAct, quantity: actionRec.quantity, face:
    actionRec.face };
} else {
    finalAction = { action: actionRec.action, quantity: actionRec.quantity,
    face: actionRec.face };
}

// Output the chosen action
postMessage(finalAction);
};

```

Implementation Notes:

- *Complexity*: Implementation can range from moderate to very high depending on how the strategy is obtained. The pseudocode assumes a precomputed `equilibriumPolicy` (perhaps loaded as a JSON or generated by a function) that maps game states to an action distribution. Generating this policy is complex (involving heavy off-line computation using CFR, deep reinforcement learning, or other game-solving techniques) [10](#) [11](#), but integrating the finished policy into the bot is straightforward: it becomes a lookup or a fixed algorithmic procedure. The pseudocode uses a simple key for the info set; a real implementation might encode the state more rigorously (including all dice counts, maybe simplified history). Accessing a dictionary or running a small neural network to get the action is well within 200ms.
- *Data Structures*: The primary structure is the policy representation. For example, a large dictionary (or a decision tree) mapping states to actions. This could be sizable – potentially many thousands of entries – but still feasible to load in a web context if compressed. Alternatively, a neural network model (with fixed weights) can be embedded and run forward passes to decide moves. Modern JavaScript can handle a small neural net in under 200ms if optimized (and WebAssembly or WebGL could further speed it, though that might be overkill). One must ensure the policy covers all reachable states; otherwise, a backup heuristic (like the baseline strategy) is needed for unseen states.
- *Computational Cost*: Low at runtime. Once training is done, using the strategy is just a matter of table lookup or evaluating a fixed formula. The heavy lifting was offloaded to the training phase (which could have taken hours offline). The example above simply finds the recommended action and picks the highest-probability move. The bot does not search during the game – it's playing a *fixed strategy* that was optimized beforehand. This means the per-move computation is minimal (constant time for lookup).
- *Determinism*: By design, this strategy eliminates randomness in execution. In true equilibrium play, one would randomize (mix strategies) to remain unpredictable [12](#), but here we choose a deterministic proxy: always taking the equilibrium action with the highest probability (or using a deterministic tiebreak rule). This makes the bot's behavior reproducible given the same state. One clever twist: some deterministic bots use their **hidden dice as a source of pseudo-randomness** to emulate mixed strategies (e.g., if an equilibrium says “bluff 30% of the time,” the bot could bluff if the sum of its dice mod 100 < 30 – this uses private information to achieve a random-like choice without calling `Math.random`). Such techniques are still deterministic from the engine's perspective and could be employed to more faithfully follow a mixed strategy. However, even without that, a largely equilibrium-based deterministic strategy will be very strong.

Studies (like a Stanford CS224R project) have shown that a neural fictitious self-play agent in a 4-player Liar's Dice can **dominate heuristic players** and achieve extremely low exploitability [7](#) [9](#). Our deterministic equilibrium bot brings those insights into a practical implementation. The ease of implementation is offset by the difficulty of obtaining the strategy, but once in place, it provides a formidable, adaptively perfect style of play that evolves over the 2500-game tournament primarily through its inherent learned policy rather than explicit on-the-fly adjustments.

Sources: The strategies above draw on probability theory and game theory fundamentals [4](#) [1](#), human expert tips on bluffing and reading opponents [3](#) [5](#), and findings from AI research on Liar's Dice (e.g. regret-minimization and self-play learning) [12](#) [9](#). Each strategy is tailored to be deterministic and efficient for a WebWorker-based simulation, while offering a distinct approach beyond the provided Baseline, AggroBluffer, ProbabilityTuned, and MomentumAdaptive bots.

[1](#) [3](#) [5](#) Liar's dice optimal strategy

https://assets.website-files.com/66f3e6db6d9b28337c56f2a2/680417c6ab8b637476e9886a_90423846074.pdf

[2](#) tournament.js

<file:///file-5rH1MyhpflQ6w19d6FrWtY>

[4](#) liars dice - Bluff - Early game strategy (5 players) - Board & Card Games Stack Exchange

<https://boardgames.stackexchange.com/questions/3403/bluff-early-game-strategy-5-players>

[6](#) [7](#) [9](#) cs224r.stanford.edu

https://cs224r.stanford.edu/projects/pdfs/CS224R_Final_Report__7_.pdf

[8](#) [10](#) [11](#) [12](#) Liar's Dice by Self-Play. With Counterfactual Regret and Neural... | by Thomas Dybdahl Ahle |

TDS Archive | Medium

<https://medium.com/data-science/liars-dice-by-self-play-3bbed6addde0>